

STAT 422/722 Spring 2016 Homework #2

Limited Solutions

Professor Adam Kapelner

Thursday 23rd February, 2017

Problem 1

We will be investigating equivalence testing.

- (a) [easy] In the context of linear or logistic regression, if you want to prove that a predictor has a linear effect on the response (controlling for other variables), what are the null and alternative hypotheses?

$$H_0 : \beta_j = 0$$

$$H_a : \beta_j \neq 0$$

- (b) [harder] In the context of linear or logistic regression, if you want to prove that a predictor does *not* have a linear effect on the response (controlling for other variables), what are the null and alternative hypotheses?

$$H_0 : \beta_j \neq 0$$

$$H_a : \beta_j = 0$$

The null and alternative are “flipped” relative to (a).

- (c) [easy] You collect four data points

predictor	response
2.47	0.50
0.57	1.95
0.84	1.91
2.18	2.51

Test the theory in (a)

Assume $\alpha = 5\%$ for testing purposes. Note: all code can be found in the file `sol.R` in this directory. Output is equivalent in JMP. R gives us:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.2698	1.0062	2.256	0.153
x	-0.3646	0.5838	-0.624	0.596

Since the p -value is $0.596 \not< 0.05$ this means we fail to reject H_0 .

- (d) [harder] Test the theory in (b). Use $\delta = 0.5$ as a margin of practical equivalence

Assume $\alpha = 5\%$ for testing purposes. We use the same output as above. We note that a 95% CI for β is $[-.3646 \pm 2 \times 0.5838] = [-1.532, 0.803]$. Since this is not a proper subset of (i.e. within) $[-\delta, +\delta] = [-0.5, 0.5]$, we fail to reject H_0 .

- (e) [difficult] How can you get both the answer to (c) and the answer to (d) at the same time? Discuss.

Problem 2

We will be investigating dredging and multiple testing corrections. I have provided a data file for you called “xyrand.csv” located here (right click and download from the browser). This file is fully random data from a standard normal and thus there is no systematic connection between the column y and any of the x_j columns.

- (a) [easy] Run a regression of y on the x_j ’s and report R^2 . Why is this R^2 not *exactly* zero?

$R^2 = 29.2\%$ and it’s not zero due to chance capitalization of x ’s being randomly related to y .

- (b) [easy] Which variables were significant and what are their significance levels? Why should any variables be significant in the first place if they’re all just $\overset{iid}{\sim}$ random realizations?

x_{26} with a p -value of .0499, x_{27} with a p -value of .0491, x_{33} with a p -value of .0448 and x_{38} with a p -value of .0320. These four variables should not be significant in the first place since there is no real correlation. They are only “significant” due to chance.

- (c) [harder] Calculate the probability you see this many significant variables *or more* in a 50-predictor linear regression. Is your number of significant variables “expected”?

This is a binomial exercise. You have a probability of falsely rejecting H_0 if H_0 is true (our case of no linear correlation) of $\alpha = 5\%$. Thus we have:

$$\mathbb{P}(N \geq 4) = \sum_{i=4}^{50} (5\%)^i (95\%)^{50-i} = 1 - F(3) \approx 24.0\%$$

This type of calculation will not be on the exam.

- (d) [easy] Calculate both a Sidak and a Bonferroni corrected individual α that preserves 5% familywise error.

Sidak is not covered on the exam. The Bonferroni α is merely $5\%/50 = 0.01\%$.

- (e) [easy] Using the Sidak and/or the Bonferroni correction, are there any significant variables anymore? Yes/no

No.

- (f) [harder] Explain precisely what I would need to simulated in this same setup with one response and 50 predictors randomly realized to expect one significant variable if the familywise correction is employed.

The Bonferroni now controls familywise error rate. The “family” here is the set of the 50 t tests. Thus it guarantees that getting one or more Type I errors in that family is at most $\alpha = 5\%$. However, if I run the entire family 20 times. That’s 20 simulations of 50 t -tests each, I will expect one Type I error since $20 \times 5\% = 1$, the expectation calculation.

- (g) [harder] Report the overall F value and it’s corresponding significance level. Explain how it is possible that there exist t tests which are significant for some linear predictors but the F test is not.

$F = 1.2285$ which has a significance level of 17.34% i.e. not significant at $\alpha = 5\%$. The F test is testing overall usefulness of the model, not individual variables. The F test “knows” that regressions with many, many variables will chance capitalize and that degree of explanatory power is expected. It only allows models to pass which are above the expected level of chance capitalization — those are “significant” models.

Problem 3

These are some conceptual questions concerning hypothesis testing, Type I errors, Type II errors and power.

- (a) [easy] Given some predetermined level of α we have two ways of setting up hypothesis tests:

H_0 : UFOs do not exist

H_a : UFOs do exist

and the inverse:

H_0 : UFOs do exist

H_a : UFOs do not exist

Which set of hypotheses should be employed and *precisely* why?

The first set is employed. H_a is the theory you wish to prove. So you are intellectually honest and assume your theory is *not* correct a priori; this is H_0 . Then, you let the data speak for itself. If the data convinces you overwhelmingly (where “overwhelmingly” is the level of skepticism defined by your level α), then you can accept your theory, but only then!

Contrast that to the second set. Here, you have begun with the belief that your theory is correct. Your theory will remain correct unless the data convinces you otherwise. This is not an intellectually honest means of reasoning.

(b) [harder] You cannot convince your friend. In the hypothesis testing framework there are two separate reasons why he could remain not convinced. What are they?

1. His α is too low i.e. he is too skeptical. No amount of data will convince him; he will never budge from his H_0 .
2. You didn’t provide enough data — you are “underpowered”.

Parenthetically, these two reasons are not either-or. In practice, it may be a combination of both. When statisticians practice, we set up α before hand and we agreed as a community it should be 5% or 1%. Conditional on accepting this community standard, the reason then would be #2.

(c) [difficult] The Shapiro-Wilk Test of Normality is used to assess normality of a given sample of data. It is a goodness-of-fit test where

H_0 : the data generating process is normal

H_a : the data generating process is not normal.

Usually, you want to prove normality (e.g. the case of testing the residuals from a linear regression). Why does this test reward small sample sizes?

Reference (a). You’ve begun by believing what you wanted to prove! So not having data is a good thing — less chance you challenge your beliefs!

Problem 4

These questions will be about extrapolation and generalization.

- (a) [harder] Is model extrapolation and model generalizability the same concept (lecture 3, slide 6)? Discuss.

Yes. Models generalize to new observations that are within the measurement / covariate space of those observations previously seen i.e. the rows in the historical dataframe you made use of when you constructed the model. Extrapolation is predicting outside of this range. It is indeed failure to generalize.

- (b) [difficult] Provide an example of a model that you use regularly that does not generalize to the observations you use to predict with it.

Prediction models for stock returns. Non-stationarity (where f is time-dependent) is another failure to generalize.

- (c) [harder] Run lines 5–27 of the lecture 3 R demos. Which of the three models would be the worst “extrapolator” and why?

Polynomial regression. The \hat{f} is too erratic outside of the region / range of the historical data. It is less likely that the true f conditional expectation function behaves like this outside of the range. Better extrapolators are linear regression and decision trees.

- (d) [difficult] You are provided a new \mathbf{x}^* with p features in which are to guess y . How would assess extrapolation? Explain.

Check each of the p measurements and see if they are outside of the ranges of those in the historical dataframe. That’s a first pass at a solution. But this is a deep problem since it’s multidimensional.

Problem 5

These questions will be about optimal design.

- (a) [easy] In the case of a simple linear regression with $n = 20$ points where x ranges from 6 to 17, what would be the optimal design?

10 observations at $x = 6$ and 10 observations at $x = 17$.

- (b) [harder] Show that for fixed n under least squares regression that the optimal design is half the points on the minimum and half the points on the maximum.

The equation for the variance of $\hat{\beta}$ in one dimension is a fraction with $\sum (x_i - \bar{x})^2$ in the denominator. The way to minimize the fraction is to maximize the denominator. This is done by choosing half of the points at the minimum and half at the maximum.

- (c) [easy] Take the case of $n = 20$ and three continuous predictors ranging in $[0, 1]$. Use JMP to create an optimal design for the model with all factorial interactions and all polynomials up until degree 3. Take the optimal design right click and make data table. Then sort the data table first by x_1 then by x_2 and by x_3 . Are they all about

the same? Yes/no.

Yes.

- (d) [difficult] The optimal design in the previous problem is what you'd probably like to do for non-parametric linear model with lots of interactions and curves. What is the takehome message in this case?
- (e) [E.C.] Your goals are prediction and you have the choice between D-optimality and I-optimality. Which is likely better and why?

Problem 6

These questions will be about logistic regression using the Telecom Churn dataset that can be downloaded [here](#).

- (a) [easy] Give an expression for the conditional mean in a logistic regression problem with p features using the standard logistic regression assumptions.
- (b) [easy] Do the multivariable logistic regression in class with target response churn (remember to delete those 6 variables which are fully collinear). Provide an interpretation on the monthly charges coefficient.
- (c) [E.C.] If we used the `cloglog` link function and got the same coefficient, what would be the interpretation?
- (d) [easy] Predict the mean probability of churn for a senior citizen who has been with the company for 36 months and then predict whether or not this person will churn.
- (e) [E.C.] In my regression output, the χ^2 value for `senior citizen` is 14.37. Calculate the standard error from this.
- (f) [harder] Test whether or not removing *both* gender and partner makes a difference in terms of linear predictive power versus the full model from (b). You will need a table of critical values of the χ^2 distribution at $\alpha = 5\%$. See below.
- (g) [easy] Use the model from (b) and use JMP to compute the AUC and misclassification error.
- (h) [difficult] Attach a graph of the false negative proportion versus the false positive proportion. Is this more useful than an ROC curve for the case of churn?
- (i) [E.C.] Create a detection error tradeoff plot for this dataset (no need to use normal deviates for x and y axes).

p	χ_p^2 critical value
1	3.84
2	5.99
3	7.81
4	9.49
5	11.07
6	12.59
7	14.07
8	15.51
9	16.92
10	18.31
11	19.68
12	21.03
13	22.36
14	23.68
15	25.00
16	26.30
17	27.59
18	28.87
19	30.14
20	31.41

- (j) [harder] Imagine the cost ratio is 7.5:1 for the more costly mistake. What is the p_0 of the optimal model?
- (k) [easy] Why is this p_0 less than the naive value of 50%?
- (l) [harder] Create a flowchart illustration of the optimal model in (j) similar to lecture 3, slide 35.

Problem 7

These questions will be about survival regression using the NetLixx dataset that can be downloaded [here](#). The response is the “time” variable and it’s measured in days (ignore the “start” and “followtime” variables. The churn variable is 1 if there is churn.

- (a) [easy] As we saw in class, the exponential model is not a great real-world model. Nevertheless, assume the model assumptions and run a standard exponential model with all three predictors. Report overall model fit, each variable’s estimate and their significance levels.
- (b) [easy] Should dredging be considered a problem here when we’re looking at all these $H_0 : \beta_j = 0$ tests? Yes/no and why.
- (c) [easy] Interpret the value of the fitted coefficient for coupon.

- (d) [easy] Predict the mean survival time for a new female of age 30 without a coupon.
- (e) [E.C.] Calculate a 95% CI for mean survival time for the new person in (d).
- (f) [E.C.] Calculate the in-sample R^2 for the non-censored rows. How does the model do?
- (g) [E.C.] Fit a Weibull model to the same data with the same response, censoring and predictors. Does it fit better? Why / why not?

Problem 8

These questions will be about our pivot from parametric linear models to non-parametric non-linear models. To illustrate, we will use the white wine dataset that can be downloaded [here](#).

- (a) [harder] Let's do an exploratory data analysis. Do Fit y by x in JMP where y is quality and x are all variables. Then, fit polynomials of degree 2 to each of them. Which ones have significant squared terms at $\alpha = 5\%$? Make sure you do a multiple testing correction.
- (b) [harder] Let's now look at all first-order interactions. These are interactions that are between two variables. Use the fit model, highlight all variables and do macros... factorial to degree (the default degree in JMP is two indicating first-order interactions). Which interactions are significant at $\alpha = 5\%$? Make sure you do a multiple testing correction.
- (c) [easy] Interpret the coefficient on `density` \times `alcohol`.
- (d) [easy] Based on your answers to (a) and (b) would it be fair to say that the true model is non-linear? Yes/no.
- (e) [easy] Based on your answers to (a) and (b) would it be fair to say that you get better predictive power if you add in some quadratic terms and interactions? Yes/no.
- (f) [easy] If you add some quadratic terms and interactions and interactions, what are the three things you are giving up in your model?
- (g) [easy] The vanilla multivariable regression gives $R^2 = 28\%$. Now fit a model with *all* interactions. Use "full factorial" in the options. What is R^2 in this new absurd model? What are the degrees of freedom?
- (h) [E.C.] What are the significant variables now? You may have to use a t-calculator.

- (i) [harder] Prove that you overfitted.

Problem 9

We will now explore the concept of overfitting.

- (a) [easy] What are you fitting when you “overfit” and why is that not a good idea?
- (b) [difficult] What are you underfitting when you “underfit” and for what purpose is this usually done?
- (c) [harder] Explain why running a linear regression with $n = p + 1$ (where the $+1$ is for β_0 , the intercept) guarantees you overfit always. Note: we are assuming each of the p features are not perfectly collinear with any of the others.
- (d) [difficult] The illustration in lecture 4, slide 32 shows the final model \hat{f} being built from the entire dataset. Why would you do this?
- (e) [easy] Explain why using the test set for more than one model generalization error estimate causes you to overfit.
- (f) [difficult] Describe a scenario where you would want to use out-of-sample validation with a test set consisting of 90% of the entire dataframe.